Sarah Chen

Staff AI / Machine Learning Engineer

sarah.chen@aiml-research.org | +1 (415) 890-1234 | San Francisco, CA

PROFESSIONAL SUMMARY

Staff Machine Learning Engineer with 8+ years experience developing enterprise LLM pipelines, deep learning architectures, and high-throughput model inference endpoints. Strong background in PyTorch, NLP, Transformers, FastAPI, Docker, and Kubernetes on AWS.

TECHNICAL SKILLS

Machine Learning: PyTorch, TensorFlow, Transformers, Hugging Face, LLMs, NLP, Scikit-learn, Vector DBs

Engineering: Python, FastAPI, Docker, Kubernetes, Kafka, PostgreSQL, Redis, AWS, CI/CD, Git

PROFESSIONAL EXPERIENCE

Staff AI Engineer — Anthropic Partner Lab

2020 – Present (4 years)

- Architected RAG inference service using Python, PyTorch, and vector databases processing 5M embeddings daily.
- Deployed scalable GPU inference endpoints in Kubernetes (AWS EKS) with sub-50ms latency.
- Published 2 research papers on transformer fine-tuning at top AI venues.

Senior ML Engineer — Cognitive Scale AI

2016 – 2020 (4 years)

- Trained and deployed deep learning classification models serving 20M users.
- Engineered streaming data ingestion pipelines using Apache Kafka and Spark.

EDUCATION

Ph.D. in Computer Science (Machine Learning) — Stanford University, 2016